

EBERHARD-KARLS-UNIVERSITÄT TÜBINGEN
Wilhelm-Schickard-Institut für Informatik
Lehrstuhl Rechnerarchitektur

Bachelor thesis

myttile

Philipp Noel von Bachmann

Betreuer: Prof. Dr. rer. nat. Andreas Zell
Wilhelm-Schickard-Institut für Informatik

Maximus Mutschler
Wilhelm-Schickard-Institut für Informatik

Begonnen am:

Beendet am: July 27, 2020

Erklärung

Hiermit versichere ich, diese Arbeit selbstständig verfasst und nur die angegebenen Quellen benutzt zu haben.

Tübingen am July 27, 2020

Max Mustermann

Contents

1. Introduction	1
1.1. Optimization in Neural Networks	1
1.1.1. Difference to normal optimization	1
1.1.2. Empirical risk minimization	1
1.1.3. Minibatch Algorithms	2
1.1.4. Challenges in Neural Networks	3
1.1.5. Algorithms	6
1.2. Related work	8
1.2.1. Generalization gap	8
1.2.2. Areas of same loss	9
1.2.3. Cosine decay with warm restart	9
1.2.4. Ensemble methods	10
2. Methods	12
2.1. Distance Function	12
2.1.1. Motivation	12
2.1.2. Mathematical approach	14
2.1.3. Pytorch implementation	18
2.2. Configuration	19
2.2.1. Library and Training	19
2.2.2. Datasets	19
2.2.3. Networks	19
2.2.4. Network hyperparameters and training loop	21
3. Summary	23
A. Abkürzungsverzeichnis	24
Bibliography	25

1. Introduction

1.1. Optimization in Neural Networks

Optimization is a core part of the current deep neural networks, as it is the foundation of the learning process. However, optimization in the context of deep learning differs from traditional optimization in several ways. This section focuses on these differences and other current challenges of optimization in deep neural networks. Most of the sections are inspired by [7].

1.1.1. Difference to normal optimization

In traditional optimization, we usually optimize on the data we want to perform later directly. However in deep learning, we usually don't have access to the test data. Consider autonomous driving. Here, the data the self-driving agent has to act on will be generated while driving, with no chance of getting it in advance. But we can capture data of other cars and optimize on them indirectly. The hope being, that the distribution of the training set is similar to the one of the later test set, so that reducing training error will result in reducing test error.

Formally, we want to reduce the test error given by

$$J(\theta) = E_{(x,y) \sim p_{data}} L(f(x; \theta), y) \quad (1.1)$$

where L is the Loss-Function, $f(x; \theta)$ the output of the network with respect to the input x and parameters θ , and y the labels for the input. This equation is known as **risk**. However, during the training process, we only have access to the training set which is distributed according to \hat{p}_{data} rather than p_{data} . Therefore, we can only optimize

$$J(\theta) = E_{(x,y) \sim \hat{p}_{data}} L(f(x; \theta), y) \quad (1.2)$$

To overcome this issue, we use a technique called empirical risk minimization.

1.1.2. Empirical risk minimization

As we have already seen, although we cannot reduce the generalization risk of equation 1.1, as we only have the training data. To convert the problem back to a normal optimization problem, we use the expectation over the empirical distribution. This is known as **empirical risk minimization**. The arithmetic mean serves as an unbiased estimator for the true mean of \hat{p}_{data} .

$$J(\theta) = E_{(x,y) \sim \hat{p}_{data}} L(f(x; \theta), y) = \frac{1}{m} \sum_{i=1}^m L(f(x^{(i)}; \theta), y^{(i)}) \quad (1.3)$$

As the training set from the empirical distribution is restricted in size, this quickly leads to overfitting, where the network memorizes the training set. Additional measurements like Regularization have to be taken to account for that problem.

1.1.3. Minibatch Algorithms

We have seen in equation 1.3 that we use the mean over the whole training set to approximate the empirical risk. From a computational perspective, this is rather expensive. That's why we normally only use a subset of the training set for each parameter update. These subsets are called batches. Some statistical considerations justify this.

The standard error of mean is given by $\frac{\sigma}{\sqrt{n}}$, where n is the number of training examples. The root in the denominator results in that, when increasing the number of training examples, the standard error of mean will only decrease sublinear. That's why it is unattractable to use large batches of training data. The second factor is redundancy in the training set. As some examples might be quite similar to each other, the mean of a subset will not differ much from the whole training data, but requires much less computation time.

Most loss-functions allow us to divide the data into batches easily. The most common example is maximum likelihood, which is defined as:

$$\theta_{ML} = \operatorname{argmax}_{\theta} p_{model}(X; \theta) \quad (1.4)$$

$$= \operatorname{argmax}_{\theta} \prod_{i=1}^m p_{model}(x^{(i)}; \theta) \quad (1.5)$$

If we convert it to log-space, the product decomposes into a sum:

$$\theta_{ML} = \operatorname{argmax}_{\theta} \sum_{i=1}^m \log p_{model}(x^{(i)}; \theta) \quad (1.6)$$

In this space, we can easily divide the sum into batches and train on them separately. The same idea can be applied to the gradient:

$$\nabla_{\theta} J(\theta) = E_{x,y \sim \hat{p}_{data}} \nabla_{\theta} \log(p_{model}(x, y; \theta)) \quad (1.7)$$

Optimization algorithms which use the whole training set are called **batch** or **deterministic** algorithms, while deterministic is preferred, as the term batch is also used in minibatch methods. The other extreme are **stochastic** or **online** methods, where only one training example is processed at a time. In between lie the methods, where more than one training example is used, but not all. These are called **minibatch** methods and are most commonly used in machine learning. A typical example for stochastic methods is stochastic gradient descent (SGD), see section 1.1.5.

Both methods have their advantages beyond the computational perspective. While large batches offer a more accurate estimate of the gradient, small batches can add a regularization effect. They add some noise to the gradient, and therefore lead the optimization algorithm to areas, where a small inaccuracy in the gradient converges to the same point. As the point is a very stable or wide minimum, this results in a regularization effect which decreases the generalization error. This is also supported by the work of Keskar et al. [13], who argued that larger batches converge to sharp minimas, whereas small converge to flat. As we will argue in flat minima are believed to be better in generalization. Therefore, a small batch which ends up in a flat minima results in better generalization capabilities.

The effect of batch size is also sensitive to the choice of the optimization algorithm. In particular, second-order algorithms suffer from a small batch size. Hessian matrices H require

a much larger sample size to be accurate than the Jacobian. Especially when H is of large condition number, this leads to the an amplification of the preexisting errors in g .

To get an unbiased estimate of the gradient, it is important to sample the mini-batches randomly. This is a problem in particular when the training data is correlated. This may arise as autocorrelation where consecutive examples are correlated for example. The problem can be overcome by sampling the minibatches uniformly out of the training data. However, that would lead to a large computational effort every time we want to construct a batch. Fortunately, it seems sufficient to shuffle and divide the training data into batches only once.

Another property is that the gradient of minibatch algorithms like SGD follow the true gradient, as long as the training data is only used once. As each datapoint from the training set is only used once, it also follows the true data distribution p_{data} . Therefore we get an estimate of the true gradient of the generalization error. When reusing the data from \hat{p}_{data} , it no longer follows p_{data} , so the estimation becomes inaccurate.

1.1.4. Challenges in Neural Networks

In the last part, we demonstrated how optimization in deep learning differs from traditional optimization from a statistical point of view. This difference is emphasized further by other factors, in particular the non-convex loss landscape. This part will summarize some of these problems and their implications for deep learning.

Local Minima

In convex optimization problems, every local minima is guaranteed to be the global minima. Therefore finding a minimum is a sufficient enough condition to stop. In neural networks however, the loss landscape is highly non-convex. This results in a local minima possibly having higher cost value than the global minima. Proofs that these local minima exist can be constructed quite easily.

One example is the weight space symmetry. Suppose you swap out nodes i and j by swapping their incoming and outgoing weights. Then the activation in the next and all subsequent layers of the networks will stay the same, but the networks are located at different places in the loss landscape, therefore creating two local minimas. Although a large number of these local minima exist, they form no problem for optimization. As mentioned, the swap will lead to the same activation, therefore the networks will perform the same on the data and the value of these local minimas will be equal.

However, there can also exist local minima with high cost compare to the global minimum. In theory this was believed to be an issue, but in practice it seemed to cause no problem. Recently, some theoretical work support these findings. Choromanska et al. [2] applied spin glass theory from physics to neural networks. This allowed them to verify, that indeed nearly all local minima are of roughly the same cost and lead to the same test error. Furthermore, the probability of a local minima with high cost quickly diminishes as the network grows in size. One reason is that most local extrema are in fact saddle points.

Saddle Points

Saddle points are another type of local extrema where the gradient is 0. In contrast to local maxima or minima where the Hessian only has negative or positive eigenvalues, the Hessian of saddle points has both positive and negative eigenvalues.

In fact, saddle points get more common for higher dimensional space. The idea of coin flipping can be used to describe this phenomena. Suppose every eigenvalue of the Hessian is generated by flipping a coin. In a low-dimensional space, it is quite likely that all of these flips will be positive or negative, resulting in minima or maxima. The higher the dimension gets however, the more likely the Hessian is to have both positive and negative eigenvalues. Therefore, there exist more saddle points. Chromanska et al. [2] also showed that these saddle points are more likely to occur in areas of higher loss.

What problems arise from saddle points depends on the choice of the optimization algorithm. For first-order algorithms, the situation is unclear. Although the gradient might get small in regions close to saddle points and as a result could slow down training, in practice it seems like it isn't a problem for gradient descent. Especially when adding Momentum to the algorithm, it is very unlikely that the gradient becomes 0, because there is no gradient in opposing direction which would decrease the momentum.

For second-order algorithms, saddle points are clearly a problem. Newton's algorithm for example explicitly solves for point with zero gradient, and is therefore attracted to saddle points or even other extrema like maxima. This is partly resolved by the introduction of saddle-free Newton.

Vanishing gradients

Today's neural network become very deep. But with increasing depth, there arises a problem called vanishing and exploding gradient. It refers to the fact that when back-propagating the gradient, it either vanishes or explodes. More formally, consider a matrix W which is multiplied by repeatedly on a computation path. If an eigendecomposition $V \text{diag}(\lambda) V^{-1}$ exists, this results in $(V \text{diag}(\lambda) V^{-1})^t = V \text{diag}(\lambda)^t V^{-1}$. Therefore, all values are scaled according to the eigenvalues of W . If these eigenvalues are between 0 and 1, the gradient will vanish. If they are larger than 1, it will explode. This especially occurs when the network becomes increasingly deep, as the power of these eigenvalues is taken. One solution for this problem is the ResNet architecture, which will be described in detail in section ??

Poor correspondence between local and global structure

The previous sections have focused on which problems we are facing when computing a gradient or updating locally. However, the local structure can often be misleading, as it doesn't reflect the global one. Even if we are able to perform the best move locally and end up in a local minima, we are not guaranteed to be in the globally best area. Figure 1.1.4 shows an example of suboptimal initialization. As the local structure for $x < 0$ doesn't reflect the global, these initializations will not lead to the global minimum.

Another issue may be that there even is no global minimum. This happens for example with the usage of the softmax, where the weights are increased without bound even when the accuracy is very good. This occurs because the actual labels for the softmax can only approach 1 with larger values, but never actually reach it. One solution to the second phenomena would be label smoothing, where instead of a hard 0-1 coding of the classification, each of the 0 are replaced by $\frac{k}{\# \text{ of output units} - 1}$, and the label for the true classification is replaced by $1 - \frac{1}{k}$. The network is now able to resemble the labels without extreme output values.

Solutions for this problem primarily focus on choosing the right initial values at the moment.




Figure 1.1.: Example of poor correspondence between local and global structure. If $x < 0$ is initialized, gradient descent will lead us in negative direction. However, this leads away from the global optimum on the right.

Initialization strategies

In the last section, the poor correspondence between local and global structure was shown. Initializations at the wrong place on the loss landscape lead to a gradient descent in a direction away from the global optimum. But as we don't know the shape of the loss landscape in advance, no initialization can guarantee a good solution or even convergence at all. However, some strategies have become widely used.

The first and only known property of initialization is that it has to be asymmetric. This means that two units that share the same input must have different weights attached to these inputs. If this is not the case, a deterministic optimization algorithm, will update both these parameters in the same way. To break this symmetry, the weights are initialized randomly. If a gaussian or uniform distribution is used doesn't really seem to matter.

However, the range of these functions seems to matter. Wide distributions have a stronger symmetry breaking effect, but come with the risk of exploding gradients as discussed in section 1.1.4. The problem of vanishing gradients arise if the distribution is too small. The statistical viewpoint suggests a initialization to small weights nevertheless. Here, large weight initialization is seen as a large gaussian prior, which and how the units interact. As we have no reason to encourage one interaction over another, the weights should be as small as possible.

Some heuristics try to balance these different motivations. One of the most famous one is the normalized initialization from Glorot [6]. For a fully connected layer of m inputs and n outputs, the weights should be initialized according to a uniform distribution:

$$W \sim U\left(-\sqrt{\frac{6}{m+n}}, \sqrt{\frac{6}{m+n}}\right) \quad (1.8)$$

The biases are normally initialized to a constant. A value of 0 seems to behave well for most applications.

Stopping criteria

In traditional optimization, simple heuristics are used to determine the point of stopping, like a gradient of 0. However many local minima exists on the loss landscape. In addition, usually we are not able to find these local extrema, as minibatch algorithms only give an approximation of the real gradient and the learning rate may lead to large step sizes which overshoot the minimum.

What we normally see in terms of performance is an initial decrease in training loss paired with an increase in validation accuracy. But after some epochs, we enter the overfitting region, where training loss keeps decreasing but validation accuracy also decreases again. In early stopping, this is used to determine the stopping point. After each epoch we store a copy of the parameter values. Then we train the network for a fixed number of epochs. At the point when the validation error starts to rise again, we can return to the parameters before. As no additional computation time is required, this is a very efficient method from a computational perspective, but may require some extra storage. This approach is very different to traditional optimization, as gradients may still be very large.

1.1.5. Algorithms

Stochastic Gradient Descent

Stochastic Gradient Descent is one of the most popular optimization algorithm in deep learning. It is also one of the most basic ones, as it only takes the gradient at the current position into account. In contrast to normal gradient descent, SGD computes the gradient on minibatches, not on the whole dataset. Nevertheless, as long as the training data is only used once, SGD gives an unbiased estimate of the true gradient (see Section 1.1.3).

Algorithm 1 Stochastic gradient descent from [7]

Require: learning rate ε

Return: a trained neural network

- 1: initialize the network, dataset and training parameters
 - 2: **while** stopping criteria is not met **do**
 - 3: sample minibatch of m examples $x^{(1)}, \dots, x^{(n)}$
 - 4: compute gradient estimate $\hat{g} = \frac{1}{m} \nabla_{\theta} \sum_i L(f(x^{(i)}; \theta), y^{(i)})$
 - 5: apply parameter update $\theta = \theta - \varepsilon \cdot \hat{g}$
 - 6: **end while**
 - 7: **return: the trained network**
-

First, the network is initialized. Then, the training loop repeats, until the stopping criteria is met. At the beginning, a minibatch is sampled. Then, the gradients are calculated, multiplied by the learning rate ε and finally applied to update the parameter values.

Learning rate decay

One of the most important hyperparameters of SGD is the learning rate. While the optimal learning rate differs for every problem, it is important to decay the learning rate with the number of epochs. Initially, it is good to choose a large learning rate. This leads to fast learning at the beginning, and avoids the algorithm getting stuck in areas of high loss. As the number of epochs increases, it is important to shrink the learning rate. SGD only is a stochastic algorithm,

therefore its gradient is inaccurate. Even when we find a local minima with a gradient of 0 on the minibatch, the true gradient will not be 0. A low learning rate secures to get close to the minimum while not overshooting it repeatedly.

[TODO: Formula 8.12, 8.13] How the learning rate is decayed varies from algorithm. A popular decay is step decay, where the learning rate is decayed by a constant factor γ after a fixed number of epochs.

$$lr = lr_0 \cdot \gamma^{\lfloor epoch/stepsize \rfloor} \quad (1.9)$$

Another popular algorithm is cosine decay, where the learning rate is decayed continuously.

$$\epsilon_t = \epsilon_{min} + \frac{1}{2}(\epsilon_{max} - \epsilon_{min})(1 + \cos(\frac{T_{cur}}{T_0}\pi)) \quad (1.10)$$

Here ϵ_{min} and ϵ_{max} define the range of the lr, T_{cur} is the current and T_0 is the maximum number of epochs. The advantage of cosine decay is that the learning rate is decayed down to 0 in the defined interval, independent of the initial learning rate. In contrast, the smallest learning rate of step decay is dependent of the initial one.

Momentum

Momentum is a popular variation of SGD. It is used to speed up the training of SGD. The term momentum is used, as the underlying idea is the similar to the physical context. Consider a frictionless ball rolling down a hill. The ball builds up speed the further it rolls down by adding the acceleration of current gradient to the velocity. The ball will speed up, until it faces an uphill, where it will slow down again.

In context of deep learning, we use the velocity v rather than only the current gradient g to update the parameters.

$$v_{t+1} = \alpha v_t - \epsilon g \quad (1.11)$$

Here, α controls how strong the past gradient is taken into account, while ϵ denotes the current learning rate. For the ball to build up velocity, it requires a constant downhill motion. The same is true for this case. The gradient can only build up, if it points in the same direction for some consecutive updates similar to a ball, which only speeds up when rolling downhill for an amount of time. Therefore, momentum speed up the gradients of parameters, whose gradient is constant in one direction. Parameters with alternating gradients for example will only experience small updates. The different orientations of their gradient will level out. Formally, if parameter p experiences the same gradient g every time, it will reach a terminal velocity of

$$\frac{\epsilon \|g\|}{1 - \alpha} \quad (1.12)$$

This also shows that α can be used to control the speed up of the training. If ϵ is kept constant, the larger α , the faster training will become. An value of 0.9 for example would lead to a speed up factor of 10, while 0.8 would lead to a speed up of 5.

In pseudo code, SGD with Momentum looks similar to normal SGD 1.1.5. The only difference is line 5 and 6, where the velocity is updated and then used to update the parameters rather than the gradient itself.

Momentum also adds an regularization effect, because it is attracted to stable or flat minima. If the minima is too small, it won't be able to stop the momentum and therefore the SGD will move on. That's similar to a ball, which won't stay in a small hole but keep on going, if it's speed is larger enough.

Algorithm 2 Stochastic gradient descent with Momentum from [7]

Require: learning rate ε

Require: momentum parameter α

Return: a trained neural network

- 1: initialize the network, dataset and training parameters
 - 2: **while** stopping criteria is not met **do**
 - 3: sample minibatch of m examples $x^{(1)}, \dots, x^{(n)}$
 - 4: compute gradient estimate $\hat{g} = \frac{1}{m} \nabla_{\theta} \sum_i L(f(x^{(i)}; \theta), y^{(i)})$
 - 5: compute velocity update $v = \alpha v - \varepsilon \hat{g}$
 - 6: apply parameter update $\theta = \theta - v$
 - 7: **end while**
 - 8: **return:** the trained network
-

1.2. Related work

1.2.1. Generalization gap

In section 1.1.1, we saw how learning differs from normal optimization, namely that we only have access to the training set and thus have to optimize indirectly. With an increasing number of training epochs, the network starts to memorize the training data. Because the data distribution of the training data is not identical to the distribution of test data, this usually leads to a better performance of the network on the training set than on the test set. The difference between those performances is known as Generalization gap.

Strategies that are developed to deal with this problem are subsumed under the term **regularization**. One common Regularizer is the L_2 Regularizer. Its goal is, to encourage the weights to stay small. Small weights have some advantages for the generalization capabilities. First of all, small weights remove the dependency of a unit to one of its inputs. Because the weights are really small, a strong activation cannot be solely achieved by the presence of one input, but rather has to rely on multiple units. Therefore small changes in the input will only cause small changes in the output, instead of the absence of one feature for example leading to a different output. This benefits the generalization, as the distribution of training and test data is slightly different. A network with L_2 regularization will nevertheless produce a quite similar output, in contrast to a normal one. Formally this can be incorporated in the loss function by adding the squared L_2 norm:

$$L = L(f(x; \theta), y) + \lambda \cdot \|\theta\|_2^2 \quad (1.13)$$

The parameter λ controls the strength of the L_2 and has to be adjusted for each problem.

Other work has gone into understanding the connection to the loss surface. Hochreiter & Schmidhuber [10] argued that flat minimas have better generalization capabilities. A flat minima is a region where the loss stays constant in contrast to sharp one, where small steps can increase the loss significantly. Therefore, flat minimas will perform constantly even for small changes in the input, whereas sharp minimas will lead to an increase in generalization error. Support also comes from the minimum description length theory, which states that fewer bits are needed to describe a flat minimum than a sharp. Lower complexity leads to a better generalization error. The idea is, that in the network we try to compress the data. The more we compress the data while also being able to resemble it, the more of the structure of the data we uncover. Therefore, a model with lower complexity can fit the underlying data better and

achieve a better generalization error. Keskar et al. [13] draw the same conclusion. They also provide a solution for finding flat minimas in using small batch sizes, see chapter 1.1.3.

Work from Dinh et al. [3] however contradicts this view. They argue that the notion of flat is problematic in the context of deep learning, as the loss surface is highly complex. Based on previous definitions from the papers above, they construct parameter values which lie on a sharp point of the surface, but are also able to generalize well. Therefore, at least some caution is needed when arguing about flatness being a reason for generalization capabilities.

1.2.2. Areas of same loss

In section 1.1.4, we showed that the poor correspondence between local and global structure may propose a major issue for optimization. Bad initializations may lead to path which moves away from the global optimum, often without chance to recover. Initialization strategies try to address this problem, but cannot guarantee to solve it.

An open question is if this suboptimal structure is present in deep neural networks. Fengxiang et al. [8] showed that there exists infinitely many local minima which are of higher cost than the global one. Furthermore, these minima are arranged in cells, where each minima of one cell is of same loss as the others and also connected with them by valleys of low loss. These cells are separated by nondifferentiable boundaries. Unfortunately, it remains unclear if these cells are of different cost. If this is the case, this would propose a major problem. When the training process gets into one of these cells, it is likely to get stuck, probably in an area of suboptimal cost. A trivial way to recover is not present at the moment.

Draxler et al. [4] get to a similar result, that local minimas are connected by valleys. In these valleys, the training loss stays the same to the one of the connected minima, while the test error rate slightly increases.

Fort & Jastrzebski [5] this in a more formal context. For each tuple of dimensions, they construct disks to describe the hyperplanes defined by them. They call these hyperplanes wedges. One property is that the valleys of low loss described above lay on these wedges. Therefore the connecting valley for two points on different wedges has to pass through the intersection of them. When viewed from a cross-section, these valleys form a tunnel. Some techniques to improve optimization have similar effects on the size of these tunnels, namely they widen them. This happens for example for higher L_2 regularization 1.2.1, smaller batch sizes 1.1.3 or higher learning rates.

1.2.3. Cosine decay with warm restart

In section 1.1.5, we introduced learning rate schedulers. The idea was to have a high learning rate at the beginning for fast improvements, and then a decrease to fine adjust the parameters. One scheduler was cosine decay, with the formula 1.10. In the paper of Loshchilov & Hutter [14], they use this scheduler in combination with another technique, called warm restart. In contrast to the naive approach, where the learning rate is only decayed once, warm restart decays until a fixed number of epochs, and then set back up to the initial learning rate. This procedure can be repeated for several times. The idea is, that in phases of high learning rate, the network explores the loss surface due to large steps. In epochs of low learning rate, the network exploits a small area to find the best parameter values. This may also add a regularization effect by letting the optimizer escape from unstable minima during high learning rate. [doubtful??]

The authors report an new state of the art result at 3.14% test error for Wide-Residual-Net 28-20. Their method also performs better when compared to step decay.

1.2.4. Ensemble methods

The general idea of ensemble methods is to combine the predictions of multiple networks to get a more accurate prediction. The fact that this leads to an improvement can be seen in a simple regression problem. Suppose there exist k models that make an error of e_i with mean 0, variance $E[e_i^2] = v$ and covariances $E[e_i e_j] = c$ on every particular prediction. If these models are combined, the variance reduces to

$$E[(\frac{1}{k} \sum_i e_i)^2] = \frac{1}{k^2} E[\sum_i (e_i^2 + \sum_{j \neq i} e_i e_j)] = \frac{1}{k} v + \frac{k-1}{k} c \quad (1.14)$$

If all models make the same predictions, so $c = v$ for all combinations of models, then the sum decomposes to v , so the prediction error is the same as before. On the other hand, if $c = 0$, the prediction error is reduced by a factor of $\frac{1}{k}$. Therefore, for ensemble methods to be successful, every model has to achieve a low prediction error, while the predictions of all models should be as different as possible. While the low error is achieved in deep learning by standard training of the models, there are several approaches to ensure that these models are different.

The first idea would be to vary the training data for the model. One way this can be realized is by k -fold cross-validation. Here, we split the training data into k different smaller datasets. Then we train k independent networks on one of these subsets. This however decreases the size of the training set for each model drastically. Therefore another common approach is bagging [1], where we draw a subset from the training data, but with replacement. Here, individual examples might occur in more than one training set. Training on different datasets leads to different models, which we can use to our advantage as we have seen above. As the number of training examples is limited however, we might not be able to sustain a sufficient low error.

To use the whole training data while also creating different models, we can alternatively alter the model itself. A naive approach would be to just use different model architectures. If we want to use the same architecture for all models, we can vary the parameter initializations. This is often enough to create models that have different predictions. However, for every initialization a network has to be trained from beginning, which can become very costly. A novel approach is to train one network, but to take snapshots of the network parameters at different steps. This approach adds no additional cost, as we only train one network. To ensure different networks, Huang et al. [12] use the method of cosine decay with warm restarts [14], as explained in section 1.2.3. Recall that at each restart, the learning rate is set up to the initial learning rate. This results in the optimizer taking larger update steps and consequently, the network parameter values will distance from their current state. The snapshots of the networks are taken before each restart, as the network converges to area of low loss when the learning rate is low. With this method, we are able to get different models with low error and no additional cost.

After we have ensured that the conditions for the models are met, we have to think about how to combine these predictions. For the case of classification with the use of softmax layer, we can use a technique called model averaging. Here, we sum the predictions of the individual models, and take the class with the highest prediction, as in standard classification. Formally, if the probability output of a model i for a given class c is p_{c_i} , then we sum the probabilities of each individual model: $p_c = \sum_i p_{c_i}$. To predict the class we take $\text{argmax}_c(p_c)$. If we have

reasons to believe in a better prediction of one model over another, we can add weights w_i to the probabilities of the individual models: $p_c = \sum_i w_i \cdot p_{c_i}$.

2. Methods

2.1. Distance Function

2.1.1. Motivation

In section 1.2.2, we got a brief overview over the loss landscape of deep neural networks and the resulting challenges for optimization. Although there exists a large number of local minima, most of them have low cost. Furthermore, they are arranged in cells, where each minima has low cost and is connected to the others via valleys of low loss.

These cells create a challenge for optimization. Suppose the learning gets into the area of one of these cells. As mentioned in chapter 1.2.2 the cells are surrounded by nondifferential boundaries and areas of higher loss. [TODO: where is this stated?] This will likely cause the algorithm to get stuck into this cell. As all of the connected minimas in this cell are of approximately same loss, there will be a boundary until the algorithm can improve, which may be higher than the global optimum. This imposes the question, if continuing to train is useful, as after one of the minima of the cell is found, the nearby minima it can reach will offer no significant improvement and other cells are out of reach.

If we reuse the ball analogy from section 1.1.5, the beginning of the training process would probably place the ball at the top of a mountain landscape. The initial training let's the ball move into one of the valleys. This valley may connect the ball to other points of low altitude, like minima are connected in a cell. When the ball is only allowed to move downhill however, the ball can never reach a spot which he is seperated by a hill. Thus, his minimum altitude is bound to the lowest place in his current valley. If we add momentum, we have seen that the ball is able to get over hills of certain size by using his momentum. Warm restarts add the possibility to jump over a hill with a high learning rate at each restart. This can lead to an escape of the cell, but without guarantee. What is more likely to happen is that after a large step, the ball is in an area of higher altitude, which will let him roll down to the same valley again in the next step.

In contrast, we try a different approach. Rather than letting the ball make one large steps and then allwoing him to move on freely, we want to constantly push the ball away from it's current position. The idea beeing if the ball is only surrounded by hills, we have to push it up one of these until it reaches the top and can then roll down into another valley, therefore escaping the cell.

Transferred back to a real network, this translates to repeatedly updating the parameter values in a way, that the new values will increase their distance to the values we want to get away from.

In Pseudo-Code, this looks the following:

First, we train the network the traiditional way - we compute the foward and backward pass of the training data with the resulting gradient, and then update the parameter values with an optimizer. After we have reached a sufficiently low error, we create the point we want to push away from, which we call **checkpoint**. The difference in the next training loop in contrast to the standard one at the beginning is how the parameter values are computed. Here, our goal is to distance from the checkpoint. This will also result in a gradient, which will be applied the

Algorithm 3 Network training with distancing**Require:** a neural network architecture and a dataset**Return:** a trained neural network

- 1: initialize the network, dataset and training parameters
- 2: **for** $i \leftarrow 1$ **to** desired number of epochs **do**
- 3: compute forward and backward pass of training data
- 4: update parameter values with optimizer
- 5: **end for**
- 6: create checkpoint we want to distance from
- 7: **for** $i \leftarrow$ next epoch **to** end **do**
- 8: compute new parameter values different to checkpoint
- 9: update parameter values with optimizer
- 10: **end for**
- 11: **return: the trained network**

same way as above.

How can we encourage a network to distance from a checkpoint, while at the same time not sacrificing the networks performance. Back in the ball analogy, instead of pushing the ball up a hill, we could place a hill at the current position, and let the ball roll down. Figure 2.1.1 shows a graphical illustration for the 2D case. An advantage of this technique is that the hill will only change the loss landscape locally. So when we have distanced from the hill, we follow the normal loss landscape again.




Figure 2.1.: Example of a function with a local minimum at $x = \frac{2}{3}$. If we place an Gaussian function on top, the minimum disappears. Note, that the black function approaches the red if we distance from the local minimum.

Formally, this hill is realised by first computing the distance of the parameter values to the

checkpoint, and then adding a penalty for small distances. If we use an exponential Kernel, the idea of the hill can be taken quite visually. Section 2.1.2 describes the mathematical approach in detail.

What happens when our penalty hill is not large enough to get the ball over the hill? This means we would get back to the present parameter values. That's not necessarily a problem, as it would mean that our current position is very robust, which would lead to a good generalization capability. Therefore our network would stay in areas with stable minima, which is a desired property.

For ensemble methods, there are other possible benefits of this approach. Recall that for an ensemble to perform well, the errors of the participating networks should be uncorellated. SGD with warm restart [14] tried to ensure this with warm restarts, where a high learning rate increases the distance of the networks and therefore leads to uncorellated errors. We could use our method as an alternative approach, where we ensure the distance between the networks not only by high learning rate but by explicitly increasing the distance. Of course, the underlying idea for both is that networks which are more distant in parameter space will behave more differently, and therefore have more uncorellated errors.

2.1.2. Mathematical approach

Distance function

As we want to measure the distance between two points, we need to define a distance function. A common choice is the euclidean distance also known as L_2 norm, which is defined as:

$$\|x\|_2 = \sqrt{\sum_i |x_i|^2} \quad (2.1)$$

This norm can also be squared to get rid of the root. Squaring does not change the direction of the gradient, and is therefore possible. To measure the distance between two parameter states θ_1 and θ_2 of the networks, this results in:

$$d(\theta_1, \theta_2) = \sum_i (\theta_{1i} - \theta_{2i})^2 \quad (2.2)$$

The size and shape of the parameters θ is not important, as each parameter is only compared to another state of itself, and is combined via a sum.

Another property is that the values the distance function can take is partly dependent on size. Consider two networks θ_1, θ_2 with the same classification task, but the size of θ_1 is larger than θ_2 . [may be confusing with before] If we assume all of the parameters are distributed the same way, then θ_1 would output a larger distance than θ_2 . However it would be desirable for the functions to be in the same bound, as it would make the transfer of hyperparameters for example possible. That's why we use a function to control for the output to be in a certain bound.

If we take a look at support vector machines, they use kernels to compute the similarity between two samples. One popular choice is the radial basis function kernel, defined as:

$$k(\theta_1, \theta_2) = \exp\left(-\frac{\|\theta_1 - \theta_2\|^2}{2\sigma^2}\right) \quad (2.3)$$

where \exp is the exponential function. With the use of 2.2, we can convert this to:

$$k(\theta_1, \theta_2) = \exp\left(-\frac{d(\theta_1, \theta_2)}{2\sigma^2}\right) \quad (2.4)$$

If we plot this function in the two dimensional case (see Figure 2.1.2), we can see some nice properties. First of all, the values are now bound between 0 and 1, regardless of the size of




Figure 2.2.: Plot of the function $f(x) = \exp(-\frac{x^2}{\sigma^2})$ for different values of σ .

θ . Second, we can see as the values of the distance get larger, the function approaches 0 asymptotically. This leads to a really small gradient for extreme distances. Consequently, the Kernel will initially push the parameters away from the checkpoint, but when this is achieved, will have little influence on the loss function. How far the function encourages to distance from the checkpoint can be controlled by the parameter σ , which defines the width of the function and can be tuned as a hyperparameter. As σ gets larger, the function becomes wider. Therefore, the influence of the distance function will reach further the larger σ is.

Loss function

This section will show how the distance function from 2.1.2 is combined with the normal loss function. The state of the art loss function for image classification, which will be used for testing is the cross-entropy-loss defined as:

$$-\sum_i \delta_{yc} \log(f(x)_i) \quad (2.5)$$

Where δ is the Kronecker-Delta function defined as:

$$\delta_{xy} = \begin{cases} 1, & \text{if } x = y \\ 0 & \text{otherwise} \end{cases} \quad (2.6)$$

To account for the distance term, we just sum it with the cross-entropy-loss:

$$L = \sum_i \delta_{yi} \log(f(x)_i) + \text{distance}(\theta, \theta_c) \quad (2.7)$$

Where θ_c is the checkpoint. Note that we can do this multiple times, so we can incorporate multiple checkpoints:

$$L = \sum_i \delta_{yi} \log(f(x)_i) + \sum_c \text{distance}(\theta, \theta_c) \quad (2.8)$$

When computing the derivative for the backpropagation, the sum decomposes into two terms, so the cross-entropy-loss will be computed the same as before. Another property we want to control for is the influence of the distance versus the cross-entropy-loss. When training is in later stages, the cross-entropy-loss may be very small. If the values of the distance function are too large in comparison, this would cause the parameters to be updated only based on the distance, which is undesirable as the performance wouldn't be taken into account anymore. The same is also true the other way around, if the distance function is too small, it wouldn't affect training at all. That's why we introduce a hyperparameter s called strength to control this:

$$L = \sum_c \delta_{yc} \log(f(x)_c) + s \cdot \text{distance}(\theta, \theta_c) \quad (2.9)$$

Effect on Gradient

In section 2.1.2, we showed how the new loss function is composed. The normal cross-entropy-loss and the distance function are added together. Therefore, taking the derivative of equation 2.7 results in

$$\nabla_{\theta} L = \nabla_{\theta} \left(\sum_i \delta_{yi} \log(f(x)_i) + \text{distance}(\theta, \theta_c) \right) \quad (2.10)$$

$$= \nabla_{\theta} \left(\sum_i \delta_{yi} \log(f(x)_i) \right) + \nabla_{\theta} \text{distance}(\theta, \theta_c) \quad (2.11)$$

The derivative of the sum decomposes in the sum of the derivatives. Therefore, the cross-entropy-loss will be added to the gradient the same as before. The new term is the distance function on the right.

First of all, it influences the direction of the gradient, so that the network parameters distance from their checkpoint values. If we take one parameter θ_i and compute the derivative of the

distance function with respect to the checkpoint θ_{c_i} , we end up with:

$$\nabla_{\theta_i} \text{distance}(\theta_{c_i}, \theta_i) = \nabla_{\theta_i} \exp\left(-\frac{\|\theta_{c_i} - \theta_i\|^2}{2\sigma^2}\right) \quad (2.12)$$

$$= \nabla_{\theta_i} \exp\left(-\frac{(\theta_{c_i} - \theta_i)^2}{2\sigma^2}\right) \quad (2.13)$$

$$= \nabla_{\theta_i} \left(-\frac{(\theta_{c_i} - \theta_i)^2}{2\sigma^2}\right) \cdot \exp\left(-\frac{(\theta_{c_i} - \theta_i)^2}{2\sigma^2}\right) \quad (2.14)$$

$$= -\frac{1}{2\sigma^2} \cdot \nabla_{\theta_i} (\theta_{c_i} - \theta_i)^2 \cdot \exp\left(-\frac{(\theta_{c_i} - \theta_i)^2}{2\sigma^2}\right) \quad (2.15)$$

$$= -\frac{1}{2\sigma^2} \cdot 2(\theta_{c_i} - \theta_i) \cdot \nabla_{\theta_i} (\theta_{c_i} - \theta_i) \cdot \exp\left(-\frac{(\theta_{c_i} - \theta_i)^2}{2\sigma^2}\right) \quad (2.16)$$

$$= -\frac{1}{2\sigma^2} \cdot 2(\theta_{c_i} - \theta_i) \cdot (-1) \cdot \exp\left(-\frac{(\theta_{c_i} - \theta_i)^2}{2\sigma^2}\right) \quad (2.17)$$

$$= \frac{1}{\sigma^2} \cdot (\theta_{c_i} - \theta_i) \cdot \exp\left(-\frac{(\theta_{c_i} - \theta_i)^2}{2\sigma^2}\right) \quad (2.18)$$

$$(2.19)$$

Only the factor $(\theta_{c_i} - \theta_i)$ will affect the direction of the gradient, the other terms will be ≥ 0 for $\forall \theta_i, \theta_{c_i} \in \mathbb{R}$ and only scale the size. Therefore, they can be disregarded and there are three cases:

- $\theta_{c_i} > \theta_i$
It follows that $\nabla_{\theta_i} \text{distance}(\theta_{c_i}, \theta_i)$ will be positive. As we apply the negative gradient to update the parameters, θ_i will get smaller and $(\theta_{c_i} - \theta_i)^2$ will be increase.
- $\theta_{c_i} < \theta_i$
 $\nabla_{\theta_i} \text{distance}(\theta_{c_i}, \theta_i)$ will be negative and the distance $(\theta_{c_i} - \theta_i)^2$ will increase the same as above.
- $\theta_{c_i} = \theta_i$
 $\nabla_{\theta_i} \text{distance}(\theta_{c_i}, \theta_i)$ will be 0. Therefore, the parameter value θ_i will stay the same.

For $\theta_{c_i} \neq \theta_i$, the gradient of the distance function will successfully increase the distance. For $\theta_{c_i} = \theta_i$ however, it will introduce no gradient at all. This would be problematic, as when a checkpoint is created, before the first update this case is true. However, the update is also composed of the cross-entropy-loss, which will likely introduce a non-zero gradient. In the next iteration, $\theta_{c_i} \neq \theta_i$ will then be true and the distance function works as expected. If Momentum is used, there is another term that prevents the gradient from becoming 0.

Besides the influence of the distance function on the direction of the gradient, we expect an influence size of the gradient. Assume that the direction of the gradient is stable across multiple iterations, which is likely if we use Momentum. After we create a checkpoint, the normal loss function would therefore increase the distance to the checkpoint by itself. As we have seen above, the gradient of the distance function will also increase the distance, therefore both gradients will point in the same direction. For a stable gradient direction of the normal loss function, the distance function therefore increases the gradient size rather than change the gradient orientation. Only when the direction for a parameter changes, the distance term will work against it.

Computational analysis

Although the network performance is one of the most important aspects of a neural network, other factors like training time have to be taken into account. Here, we will analyze how the distance term affects training time. Algorithm 4 shows that at the beginning, the network is trained as usual. Therefore, there is no additional training time in this stage. After we add the checkpoint however, the distance function is added to the loss. In section 2.1.2 we have seen that there is no interaction between the gradient of the cross-entropy-loss and the distance function. The same is true for the forward pass. Therefore, the only added computation time raises from the distance term.

The number of parameters of the network stays constant over the epochs. Therefore, the number of mathematical operations in the forward and backward pass of the distance term stays constant. Let's assume the comparison time for two numbers is also constant regardless of their size. Then it follows, that the distance function creates a constant additional amount of time, denoted by d . As we have to do this operation every n times we compute our loss, we add $O(n \cdot d)$ to our cost $O(n \cdot l + n \cdot d)$. As $O(n \cdot l + n \cdot d) = O(n \cdot (l + d)) = O(n)$, the complexity of the algorithm stays the same. However, we have a constant for the checkpoint. If we add more checkpoints, there are again no interactions. Therefore, every new checkpoint should add the same additional cost. The actual value of the additional cost will be discussed in the results, section ??.

Multiple Checkpoints

2.1.3. Pytorch implementation

Checkpoint creation

To measure the distance, we have to create the checkpoint. The model parameters in pytorch are stored as matrices for each layer and can be accessed via `model.parameters()`, which outputs an iterable. We therefore opt to keep this structure and save the parameters in a list.

Algorithm 4 Checkpoint

```
import torch

def create_checkpoint(model):
    checkpoint = []
    for param in model.parameters():
        newparams = param.clone().detach().to(device='cuda')
        checkpoint.append(newparams)

    return checkpoint
```

First, the checkpoint list is initialized. Then we iterate over the model parameters. For each layer, we have to clone the parameters in order to create new variables, and not just pointers to the existing ones. In addition they have to be detached, to remove connection of the gradient to the model parameters.

Distance function

The L2 norm is implemented the following way: We iterate over the checkpoint and the current parameter values. For each layer, we compute the difference, and then square and add the values to our distance.

Algorithm 5 L2 norm

```
import torch

def computedistance(checkpoint, model):
    distance = 0
    for checkpoint_param, param in zip(checkpoint, model.parameters()):
        difference = torch.add(-param, checkpoint_param)
        distance += torch.sum(torch.mul(difference, difference))
    return distance
```

2.2. Configuration

2.2.1. Library and Training

The code for the network was written in Python [Version], with the use of Pytorch [15] as the machine learning library. Data that occurred during training was logged and plotted with Tensorboard. For a detailed overview of the code, see Appendix [add]. Training was done on the TCML Cluster of the university of Tübingen.

2.2.2. Datasets

Two common datasets were used.

- **MNIST**
MNIST is a classical example of hand-written digits from 0 to 9 that have to be classified accordingly. It contains a train set of 60000 and a test set of 10000 examples.
- **CIFAR-10** This dataset consists of 60000 32x32 colour images with 10 different classes. The dataset is separated into 50000 train and 10000 test images. The state of the art accuracy is 97.3%, reported from Kolesnikov et al. (2019). The dataset can be accessed via <https://www.cs.toronto.edu/~kriz/cifar.html>.

2.2.3. Networks

ResNet

The Residual Network (ResNet) architecture was first proposed by He et al. [9]. The idea of the ResNet was, that it is more easy for a network to learn the residuals rather than the full transformation of an input. Formally, consider the input x . The network transforms x according to $H(x)$. Rather than letting the network do the full transformation, ResNet produces $H(x) = x + f(x)$, where $f(x)$ is the residual transformation learned by the network, and x the identity data which is realized by a skip connection. Figure 2.2.3 shows a graphical illustration.

Beside the assumably easier to learn representation, ResNet also solves other problems. The issue of the vanishing gradient as discussed in section 1.1.4 for example is tackled, as skip connections backpropagate the gradient better to earlier layers of the network. This allows for deeper networks.




Figure 2.3.: Residual Block from [9, Page 2]

ResNet creates basic building blocks by applying a convolution, a Relu and another convolution before adding the skip connection with a final Relu. Figure 2.2.3 illustrates this block. The blocks are then stacked after each other. To let the network learn a more complex representation, after a number of block, they increase the number of channels. Here, the identity mapping of the skip connection is replaced by a pointwise convolution to increase the number of channels. Alternatively, the new channels are padded with 0.

MobileNetV2

MobileNet was first introduced by Howard et. al [11] as a lightweight neural network for the use on mobile devices. To reduce the computation effort of the network, they made use of the depth-wise separable convolution. Consider a $32 \times 32 \times 3$ image of the CIFAR-10 dataset. A traditional 3×3 convolution with a stride and padding of 1 would produce an output of $\frac{32+2-3}{1} + 1$ the shape $32 \times 32 \times 1$. To get more channels, we would need more kernels. For a total of k output channels, we would need to do $3 \times 3 \times 3 \times 32 \times 32 \times k$ multiplications. Instead of doing the computation in one kernel, depth-wise separable convolution divides it into two kernels. First they perform a depthwise convolution, where a kernel of $3 \times 3 \times 1$ is applied to every channel of the input, so in this case we end up same size of $32 \times 32 \times 3$. To get to k channels, they use a pointwise convolution across the channels. This is a $1 \times 1 \times 3$ convolution, which upscales the image to more channels. So if we want k output channels, we need k $1 \times 1 \times 3$ kernels. The benefit of this technique is that, while getting the same output size, we only need $32 \times 32 \times 3 \times 3 \times 3$ multiplications in the first step, and $32 \times 32 \times 1 \times 1 \times 3 \times k$ in the second step. This is way less than the from the beginning. Figure 2.2.3 shows a graphical illustration of this idea.

With the second iteration called MobileNetV2 [16], they added the concept of inverted residuals or as they call it bottlenecks. The idea is, that intermediate representations should be low rather than high dimensional. The bottleneck layer starts with a matrix with a small number of depth channels. In the first step, the representation is expanded by a pointwise convolution, followed by a Relu6. The expansion factor here controls how much channels the intermediate representation will have. Then a depthwise convolution is applied as a transformation step,




Figure 2.4.: Depthwise Separable Convolution from [16, Page 3]

On the left, a regular convolution with a $3 \times 3 \times \text{Depth}$ kernel is drawn. On the right, this convolution is separated in depthwise convolution with a $3 \times 3 \times 1$ Kernel, followed by a pointwise convolution with a $1 \times 1 \times \text{Depth}$ Kernel. This method saves computation time over the standard method, while also allowing interactions between channels.

followed again by a Relu6. Finally, a pointwise convolution is applied again, but this time without a nonlinear transformation following afterwards. The idea is that when compressing the data back to a low dimensional space, a nonlinearity would only lead to a loss of information. Finally, a skip connection from the input to the output is applied to get a residual mapping from input to output.

For ImageNet, MobileNetV2 reaches a top-1 accuracy of 72%, which outperforms the competitors like MobileNetV1 or ShuffleNet 1.5 while having a similar number of parameters.




Figure 2.5.: Inverted Residual Block from [16, Page 3]

First, a Relu 6 followed by a Pointwise Convolution is applied to map the data to a higher dimensional space. Then a depthwise convolution is applied, followed again by a Relu6. Finally, the data is compressed back to the original size. This time, there is no Relu applied to avoid the loss of information.

2.2.4. Network hyperparameters and training loop

For both networks, the base hyperparameters are the same. Note that these parameters are only standard configurations and may be changed to investigate the effects:

- Optimization Algorithm
Stochastic Gradient descent combined with Momentum is used. A λ of 0.9 is set for Momentum.
- Learning rate
Initially $1e-2$

- Batch size
128
- Regularization
An L_2 Regularization as described in chapter 1.2.1 is used with a factor of $1e-3$.
- Loss Function
Cross-Entropy Loss
- Number of epochs
The number of epochs varies. As this is an explorative analysis, the training is often run for a large number of epochs, to investigate long term changes. Usually an epoch number of 900 is used.

3. Summary

A. Abkürzungsverzeichnis

tRNA Transfer-RNA

Bibliography

- [1] L. Breiman. Bagging predictors. *Machine learning*, 24(2):123–140, 1996.
- [2] A. Choromanska, M. Henaff, M. Mathieu, G. B. Arous, and Y. LeCun. The loss surfaces of multilayer networks. In *Artificial intelligence and statistics*, pages 192–204, 2015.
- [3] L. Dinh, R. Pascanu, S. Bengio, and Y. Bengio. Sharp minima can generalize for deep nets. *arXiv preprint arXiv:1703.04933*, 2017.
- [4] F. Draxler, K. Veschgini, M. Salmhofer, and F. A. Hamprecht. Essentially no barriers in neural network energy landscape. *arXiv preprint arXiv:1803.00885*, 2018.
- [5] S. Fort and S. Jastrzebski. Large scale structure of neural network loss landscapes. In *Advances in Neural Information Processing Systems*, pages 6709–6717, 2019.
- [6] X. Glorot and Y. Bengio. Understanding the difficulty of training deep feedforward neural networks. In *Proceedings of the thirteenth international conference on artificial intelligence and statistics*, pages 249–256, 2010.
- [7] I. Goodfellow, Y. Bengio, and A. Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [8] F. He, B. Wang, and D. Tao. Piecewise linear activations substantially shape the loss surfaces of neural networks. *arXiv preprint arXiv:2003.12236*, 2020.
- [9] K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.
- [10] S. Hochreiter and J. Schmidhuber. Flat minima. *Neural Computation*, 9(1):1–42, 1997.
- [11] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. *arXiv preprint arXiv:1704.04861*, 2017.
- [12] G. Huang, Y. Li, G. Pleiss, Z. Liu, J. E. Hopcroft, and K. Q. Weinberger. Snapshot ensembles: Train 1, get m for free. *arXiv preprint arXiv:1704.00109*, 2017.
- [13] N. S. Keskar, D. Mudigere, J. Nocedal, M. Smelyanskiy, and P. T. P. Tang. On large-batch training for deep learning: Generalization gap and sharp minima. *arXiv preprint arXiv:1609.04836*, 2016.
- [14] I. Loshchilov and F. Hutter. Sgdr: Stochastic gradient descent with warm restarts. *arXiv preprint arXiv:1608.03983*, 2016.

- [15] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Kopf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala. Pytorch: An imperative style, high-performance deep learning library. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems 32*, pages 8024–8035. Curran Associates, Inc., 2019.
- [16] M. Sandler, A. Howard, M. Zhu, A. Zhmoginov, and L.-C. Chen. Mobilenetv2: Inverted residuals and linear bottlenecks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 4510–4520, 2018.